

Arm® Ethos™-N NPU Static Performance Analyzer Tool

This confidential and proprietary software may be used only as authorised by a licensing agreement from Arm Limited

(C) COPYRIGHT 2018-2020 Arm Limited

ALL RIGHTS RESERVED

The entire notice above must be reproduced on all authorised copies and copies may only be made to the extent permitted by a licensing agreement from Arm Limited.

The **Arm® Ethos™-N NPU Static Performance Analyzer Tool** (SPA) provides a performance report using Arm NN and parts of the Ethos-N NPU driver stack, without needing to execute the inference. This can be done without accessing the hardware, Ethos-N Kernel Driver or Ethos-N Driver Library.

The Static Performance Analyzer consists of two parts:

- The `AnalyzeNetwork` tool, shipped as part of the SPA package. See the `README.md` for instructions on how to build it.
- **This** Jupyter notebook, showing a report based on data generated by running `AnalyzeNetwork`.

All values unless stated are **predicted future** values based on future software improvements scheduled. For current performance, please set the `CURRENT` flag to 1.

Please note that there are some limitations of the performance figures reported here. See the `Limitations` section in the `README.md` file included in the distribution package.

In [1]:

```
# For filtering operator names
import re
# For data manipulation within this notebook
import pandas
# Setup pandas global parameters to help visualisation
pandas.set_option('display.max_colwidth', 120)
pandas.set_option('display.max_rows', None)
pandas.set_option('display.max_columns', None)
# For presentation of tables within this notebooks
from IPython.display import display
# For presentation of graphs within this notebooks
%matplotlib inline
# For presentation of SVG graphs within this notebooks
from IPython.display import HTML
# For the presentation of the notebook itself, use the full screen width
display(HTML(''
<style>
    div#notebook-container    { width: 95%; }

```

```

div#menubar-container      { width: 65%; }
div#maintoolbar-container  { width: 99%; }
</style>
''' )

```

```

# The spa module is required to extract the data from the raw json report
import spa
# Apply the global style for all charts
spa.apply_global_style()

```

Notebook Input

Input Required:

The following parameters need to be set by the user in order to run this notebook successfully

OUTPUT_DIR The PERFORMANCE_OUTPUT_DIR value specified in the config file
This is the path to the output from executing the driver in Performance Mode

CLK The Ethos-N NPU Clock Frequency to use, in Hertz

BANDWIDTH The DRAM Bandwidth of the system to use, in Bytes per second

In [2]:

```

# Define the output folder path generated from your run
OUTPUT_DIR = './perfreport_output/'

# Define the Clock Frequency of the Ethos-N in the target system - in Hertz
CLK = 1.25e9 # 1.25GHz

# Define the DRAM Bandwidth of the target system - in Bytes per second
BANDWIDTH = 16e9 # 16GBps

```

Summary

The below table shows a summary view of some key metrics.

In [3]:

```

# Parse the given json file
results = spa.parse_output_dir(OUTPUT_DIR)

# Get the parameters of Subgraph 0 (assuming they are the same for the rest of
params, _, _ = results[0]

summary_table = pandas.DataFrame()
for subgraph, (_, data, _) in enumerate(results):
    summary = spa.get_summary(data, BANDWIDTH, CLK)
    summary.name = 'Subgraph_{}'.format(subgraph)
    summary_table = pandas.concat([summary_table, summary], axis=1, sort=False)

summary_table = summary_table.style.set_caption(params['Variant']).set_table_
    'selector': 'caption',
    'props': [
        ('color', 'black'),

```

```

        ('font-size', '20px'),
        ('text-align', 'center')
    ]
}]
display(summary_table)

```

 Arm® Ethos™-N NPU Static Performance Analyzer Tool

Parsing file:
 ./perfreport_output/subgraph_0/report.json
 Config:
 Variant Ethos-N78_4TOPS_4PLE_RATIO
 SramSizeBytesOverride 0
 ActivationCompressionSavings 0
 WeightCompressionSavings Not Specified
 Current 0

Ethos-N78_4TOPS_4PLE_RATIO

	Subgraph_0
Total Cycle Count	8373024.000000
Time Taken (ms)	6.700000
Inferences per second	149.290000
Network Complexity (Millions of Operations)	54262.240000
Total DRAM Transactions (MB)	31.330000
Total SRAM Transactions (MB)	37.160000
Total Input Transactions (MB)	42.270000
Total Output Transactions (MB)	23.260000
Total Weights Transactions (MB)	2.960000
Average weight compression	0.620000

Operations

List of supported operations

This is the list of **unique** operators found in the report that are supported.

```

In [4]: regex = re.compile(r'(:[0-9+]{2}')
ops = []
for (_, data, _) in results:
    ops.extend(regex.sub('', y) for x in data['Operation', 'Names'] for y in
    uniqops = set(ops)
    print(*uniqops, sep='\n')

```

MaxPool2D
 Conv2D
 Conv2D:RELU6

List of unsupported operations

This is the list of operators that are **not** supported!

```
In [5]: df = pandas.Series(dtype=str)
for subgraph, (a, b, issues) in enumerate(results):
    print("a: {}".format(a))
    print("b: {}".format(b))
    s = pandas.Series(issues)
    s.index = [('Subgraph {}'.format(subgraph), idx, name) for idx, name in s]
    df = df.append(s)
if not df.empty:
    df.index = pandas.MultiIndex.from_tuples(df.index)
df.name = 'Issues'
df.to_frame()
```

```
a: {'Variant': 'Ethos-N78_4TOPS_4PLE_RATIO', 'SramSizeBytesOverride': 0, 'ActivationCompressionSavings': 0, 'WeightCompressionSavings': 'Not Specified', 'Current': 0}
```

```
b:          Input
DramParallelBytes  DramNonParallelBytes  SramBytes  NumCentralStripes
0                0.0                0.0    2359296.0             12.0
1                0.0                0.0   18874368.0              0.0
2                0.0                0.0   4915200.0             32.0
3          2162688.0          196608.0            0.0             24.0
4          2516582.0          629145.0            0.0             12.0
5                0.0                0.0   1572864.0             12.0
6          655360.0          131072.0            0.0             12.0
7          1258291.0          314572.0            0.0              6.0
8                0.0                0.0    786432.0              6.0
9                0.0                0.0    393216.0              1.0
10               0.0                0.0   1572864.0             6.0
11               0.0          786432.0            0.0              1.0
12          314572.0          78643.0            0.0              1.0
13               0.0                0.0    786432.0              1.0
14               0.0          393216.0            0.0              1.0
15               0.0                0.0    786432.0              1.0
16               0.0          786432.0            0.0              1.0
```

```
          Output
NumBoundaryStripes  NumReloads  DramParallelBytes  DramNonParallelBytes
0                0.0            0.0    1887436.0    471859.0
1                0.0            0.0    4669440.0     49152.0
2                16.0            0.0    1887436.0    471859.0
3                0.0            0.0    3014656.0    131072.0
4                0.0            0.0            0.0            0.0
5                0.0            0.0    629145.0    157286.0
6                0.0            0.0   1441792.0    131072.0
7                0.0            0.0            0.0            0.0
8                0.0            0.0            0.0            0.0
9                0.0            0.0            0.0            0.0
10               0.0            0.0    629145.0    157286.0
11               0.0            0.0    380928.0     12288.0
12               0.0            0.0            0.0            0.0
13               0.0            0.0    314572.0     78643.0
14               0.0            0.0    749568.0     36864.0
15               0.0            0.0    629145.0    157286.0
16               0.0            0.0            0.0     57600.0
```

```
          SramBytes  NumCentralStripes  NumBoundaryStripes  NumReloads
0                0.0                12.0                0.0            0.0
1                0.0                96.0                0.0            0.0
2                0.0                32.0                0.0            0.0
3                0.0                24.0                0.0            0.0
4          1572864.0                12.0                0.0            0.0
5                0.0                12.0                0.0            0.0
6                0.0                12.0                0.0            0.0
7          786432.0                 6.0                0.0            0.0
8          393216.0                 6.0                0.0            0.0
9          1572864.0                 8.0                0.0            0.0
10               0.0                 6.0                0.0            0.0
```

11	0.0	32.0	0.0	0.0
12	786432.0	22.0	0.0	0.0
13	0.0	11.0	0.0	0.0
14	0.0	22.0	0.0	0.0
15	0.0	64.0	0.0	0.0
16	0.0	1.0	0.0	0.0

Weights

	DramParallelBytes	DramNonParallelBytes	SramBytes	NumCentralStripes
0	0.0	0.0	0.0	0.0
1	0.0	512.0	0.0	1.0
2	4096.0	0.0	0.0	1.0
3	0.0	8704.0	0.0	1.0
4	0.0	1792.0	0.0	1.0
5	6656.0	0.0	0.0	1.0
6	0.0	16384.0	0.0	1.0
7	0.0	4608.0	0.0	1.0
8	16384.0	0.0	0.0	1.0
9	100352.0	0.0	0.0	8.0
10	43520.0	0.0	0.0	1.0
11	174592.0	5632.0	0.0	32.0
12	591360.0	28160.0	0.0	22.0
13	118272.0	0.0	0.0	11.0
14	591360.0	28160.0	0.0	22.0
15	1196032.0	0.0	0.0	64.0
16	0.0	22528.0	0.0	1.0

Mce

	NumBoundaryStripes	NumReloads	CompressionSavings	Operations	CycleCount
0	0.0	0.0	0.000000	0.000000e+00	0.0
1	0.0	0.0	0.000000	3.774874e+07	196608.0
2	0.0	0.0	0.604938	4.076863e+09	589824.0
3	0.0	0.0	0.685185	2.717909e+09	294912.0
4	0.0	0.0	0.125000	2.013266e+08	49152.0
5	0.0	0.0	0.638889	1.811939e+09	196608.0
6	0.0	0.0	0.777778	1.811939e+09	196608.0
7	0.0	0.0	0.437500	2.013266e+08	49152.0
8	0.0	0.0	0.777778	1.811939e+09	196608.0
9	0.0	0.0	0.829861	3.623879e+09	393216.0
10	0.0	0.0	0.667969	8.053064e+08	196608.0
11	0.0	0.0	0.847222	7.247757e+09	786432.0
12	0.0	0.0	0.868707	7.247757e+09	786432.0
13	0.0	0.0	0.774414	8.053064e+08	196608.0
14	0.0	0.0	0.868707	7.247757e+09	786432.0
15	0.0	0.0	0.873264	1.449551e+10	1572864.0
16	0.0	0.0	0.706667	1.179648e+08	30720.0

Ple

Overhead

	CycleCount	NumOfPatches	Operation	StartAndEndCycleCount	StripesCycleCount
0	0.0	0.0	0.0	0.0	1200.0
1	1179648.0	98304.0	5.0	0.0	5030.0
2	442368.0	36864.0	5.0	0.0	9810.0
3	73728.0	12288.0	10.0	0.0	6770.0
4	36864.0	6144.0	10.0	0.0	3410.0
5	147456.0	12288.0	5.0	0.0	3410.0
6	36864.0	6144.0	10.0	0.0	3410.0
7	18432.0	3072.0	10.0	0.0	1730.0
8	73728.0	6144.0	5.0	0.0	1730.0
9	36864.0	6144.0	10.0	0.0	2290.0
10	18432.0	3072.0	10.0	0.0	1730.0
11	73728.0	6144.0	5.0	0.0	9010.0
12	18432.0	3072.0	10.0	0.0	6210.0
13	9216.0	1536.0	10.0	0.0	3130.0
14	18432.0	3072.0	10.0	0.0	6210.0
15	18432.0	3072.0	10.0	0.0	17970.0
16	1440.0	240.0	10.0	0.0	330.0

	Parent	Operation
StripesNonParallelCycleCount	Ids	Ids
0	0.0	[[1]]
1	0.0	[[0]] [1, 2, 3, 4]
2	0.0	[1] [5, 6, 7, 8]
3	0.0	[2] [9, 10, 11]
4	0.0	[3] [12, 13, 14]
5	0.0	[4] [15, 16, 17, 18]
6	0.0	[5] [19, 20, 21]
7	0.0	[6] [22, 23, 24]
8	0.0	[7] [25, 26, 27, 28]
9	0.0	[8] [29, 30, 31]
10	0.0	[9] [32, 33, 34]
11	0.0	[10] [35, 36, 37, 38]
12	0.0	[11] [39, 40, 41]
13	0.0	[12] [42, 43, 44]
14	0.0	[13] [45, 46, 47]
15	0.0	[14] [48, 49, 50]
16	0.0	[15] [51, 52, 53, 54]

	Names
0	[Input from input]
1	[None, None, Conv2D:0:0, MaxPool2D:0:1]
2	[None, None, Conv2D:0:2, MaxPool2D:0:3]
3	[None, None, Conv2D:0:4]
4	[None, None, Conv2D:0:5]
5	[None, None, Conv2D:0:6, MaxPool2D:0:7]
6	[None, None, Conv2D:0:8]
7	[None, None, Conv2D:0:9]
8	[None, None, Conv2D:0:10, MaxPool2D:0:11]
9	[None, None, Conv2D:0:12]
10	[None, None, Conv2D:0:13]
11	[None, None, Conv2D:0:14, MaxPool2D:0:15]
12	[None, None, Conv2D:0:16]
13	[None, None, Conv2D:0:17]
14	[None, None, Conv2D:0:18]
15	[None, None, Conv2D:0:19]
16	[None, None, Conv2D:0:20, Conv2D:0:20:RELU6]

Out[5]:

	Issues
Subgraph 1 0	NaN Could not be estimated: Please provide a mapping file entry for this operation
2	NaN Could not be estimated: Please provide a mapping file entry for this operation
3 Conv2D:0:0	Could not be estimated: Please provide a mapping file entry for this operation

Breakdown

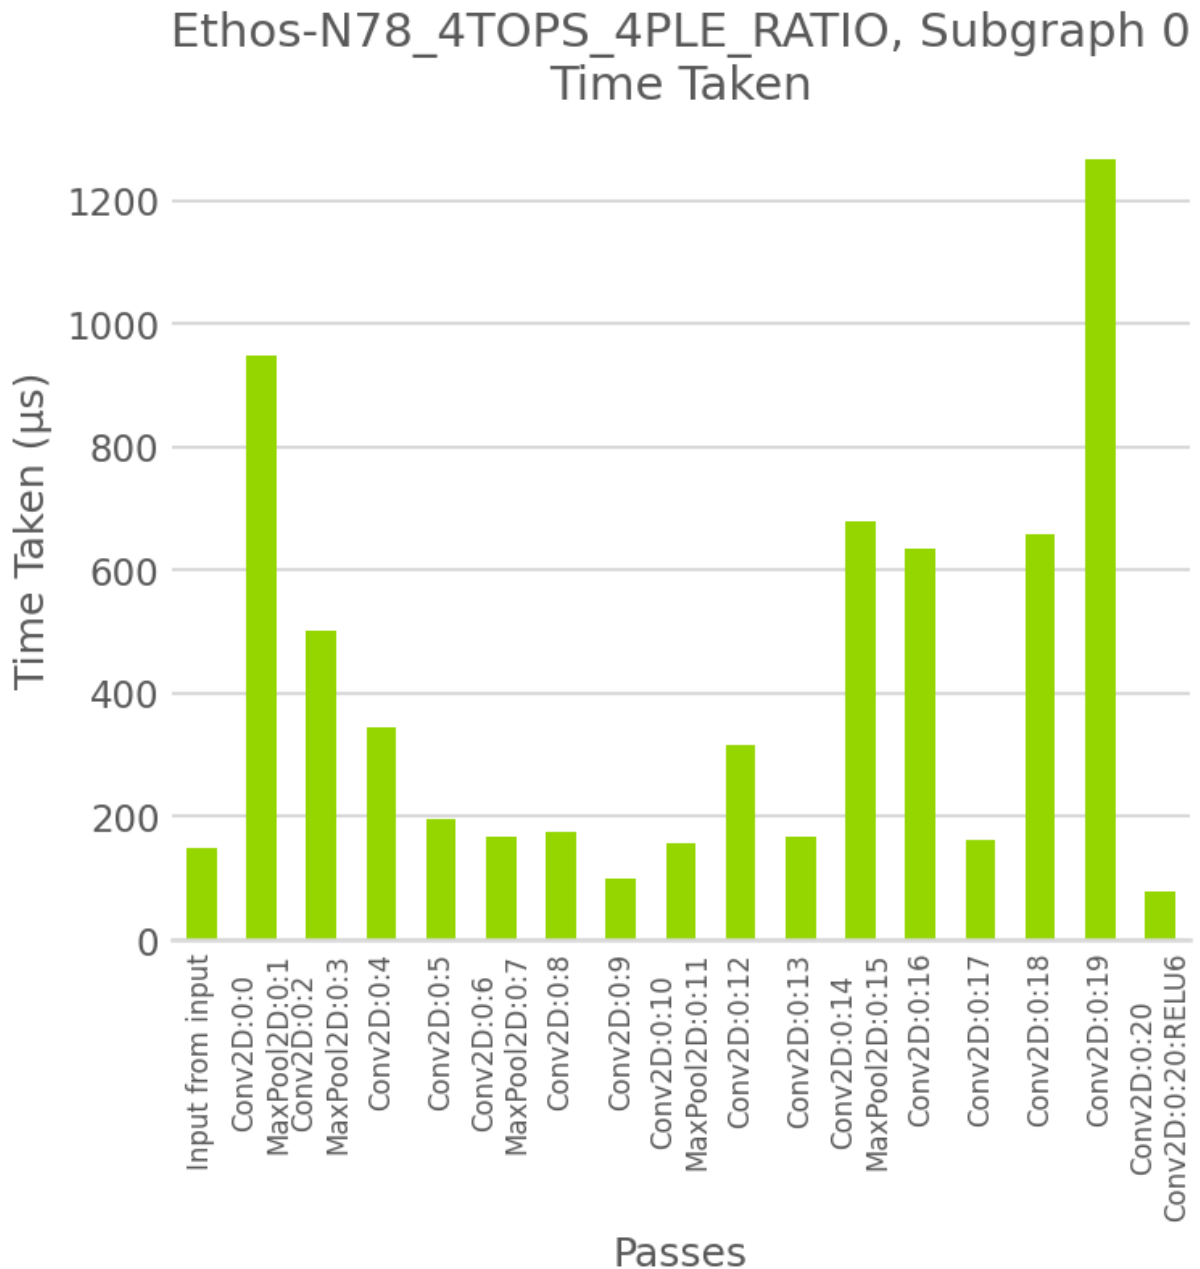
Time Taken

This is a breakdown of the time taken for each pass.

It requires a global `BANDWIDTH` value to be set to specify what the system DRAM bandwidth limit is, and it also requires a global `CLK` value to be set to specify what the clock frequency of the NPU is.

```
In [6]: for subgraph, (_, data, _) in enumerate(results):
        cycles = spa.calc_total_cycles_range(data, BANDWIDTH / CLK)['Predicted']
        # Convert Cycles to Time, and scale to microseconds
```

```
df = cycles.div(CLK).mul(1e6)
ax = df.plot.bar(color=spa.ARM_COLORS['green'])
spa.apply_formatting(ax, '{}', Subgraph {} \nTime Taken'.format(params['Var
```



DRAM Bandwidth

This is a breakdown of DRAM Bandwidth per pass, showing both Reads and Writes (stacked).


It requires a global `BANDWIDTH` value to be set to specify what the system DRAM bandwidth limit is, and it also requires a global `CLK` value to be set to specify what the clock frequency of the NPU is.

The first chart illustrates the bandwidth on a per-pass basis.

The second chart illustrates the bandwidth over time.

```
In [7]: for subgraph, (_, data, _) in enumerate(results):
df = spa.extract_dram_bandwidth(data, BANDWIDTH, CLK)
df = df[df.columns[::-1]] # Plot Writes first then Reads
ax = df.plot.bar(stacked=True, color=spa.COOL_COLOR_PALETTE[::-1])
```

```
ymax = BANDWIDTH / 1e9 # Scale to GBps
spa.apply_formatting(ax, '{}', Subgraph {} \n DRAM Bandwidth'.format(params[
```



```
for subgraph, (_, data, _) in enumerate(results): df = spa.extract_dram_bandwidth(data, BANDWIDTH, CLK)
timetaken = spa.extract_time_taken(data, BANDWIDTH, CLK) # Scale the timeline to ms timetaken *= 1e3 #
Plot a custom stacked bar chart where the width of the bar is the time taken spa.plt.bar(x=timetaken.cumsum(),
height=df['Write'], width=-timetaken, align='edge', label='Write', color=spa.ARM_COLORS['dark blue'])
spa.plt.bar(x=timetaken.cumsum(), height=df['Read'], bottom=df['Write'], width=-timetaken, align='edge',
label='Read', color=spa.ARM_COLORS['blue']) # Get the current axes for manipulating ax = spa.plt.gca() # Set
the x-axis limits ax.set_xlim(0, timetaken.sum()) ymax = BANDWIDTH / 1e9 # Scale to GBps
spa.apply_formatting(ax, '{}', Subgraph {} \n DRAM Bandwidth'.format(params['Variant'], subgraph), 'Time (ms)',
'Bandwidth (GBps)', ymax=ymax, figwidth=12)
```

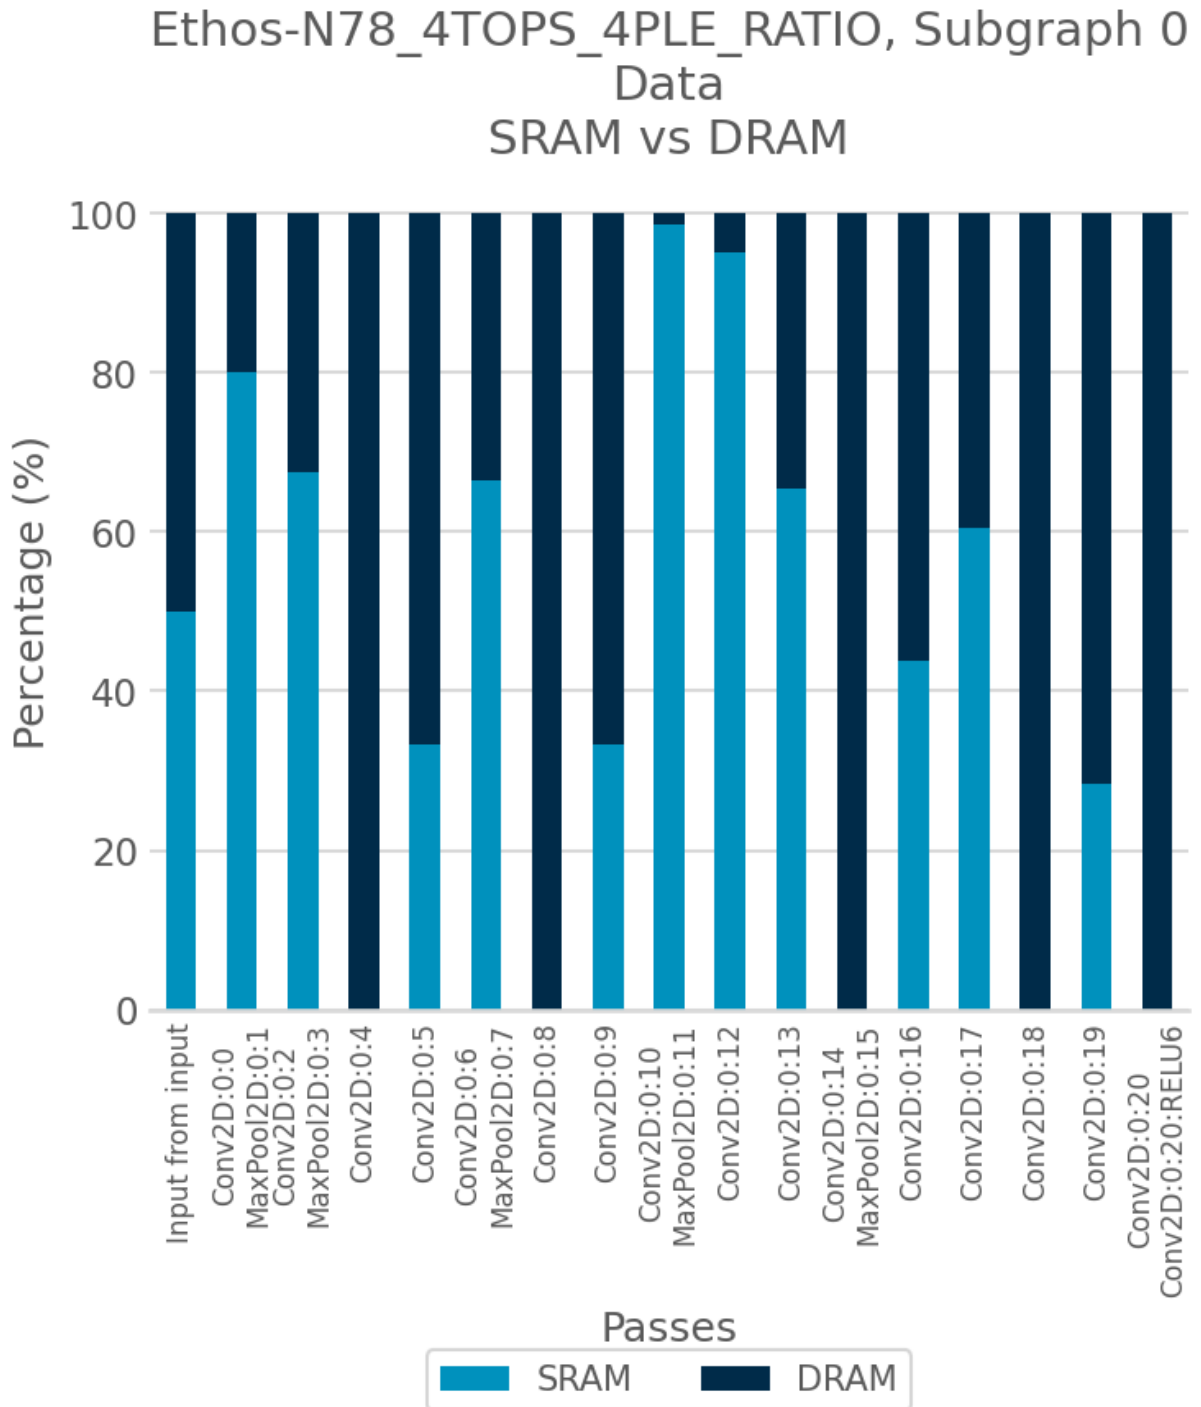
Data SRAM vs DRAM

This is a breakdown of where the data is for each pass (both Read and Write).

```
In [8]: for subgraph, (_, data, _) in enumerate(results):
```



```
df = spa.extract_sram_v_dram(data)
ax = df.plot.bar(stacked=True)
spa.apply_formatting(ax, '{}', Subgraph {} \nData \nSRAM vs DRAM'.format(par
```



DRAM Memory Transactions

This is a breakdown of DRAM memory transactions per pass.

The first chart illustrates the split between Inputs, Outputs and Weights for the DRAM transactions that occur during that pass.


The second chart illustrates the split between Reads and Writes for the same data.

The dotted horizontal line is the AVERAGE for each.

```
In [9]: for subgraph, (_, data, _) in enumerate(results):
```


```
df = spa.extract_dram_transaction_type(data)
ax = df.plot.bar(stacked=True, color=spa.COOL_COLOR_PALETTE)
spa.apply_mean(ax, df, color=spa.COOL_COLOR_PALETTE)
spa.apply_formatting(ax, '{}', Subgraph {} \nDRAM Memory Transactions \nInput
```

Ethos-N78_4TOPS_4PLE_RATIO, Subgraph 0 DRAM Memory Transactions Input vs Output vs Weights



```
In [10]: for subgraph, (_, data, _) in enumerate(results):
df = spa.extract_dram_transaction_rw(data)
ax = df.plot.bar(stacked=True)
spa.apply_mean(ax, df)
spa.apply_formatting(ax, '{}', Subgraph {} \nDRAM Memory Transactions \nReac
```

Ethos-N78_4TOPS_4PLE_RATIO, Subgraph 0 DRAM Memory Transactions Read vs Write



Weight Compression Savings

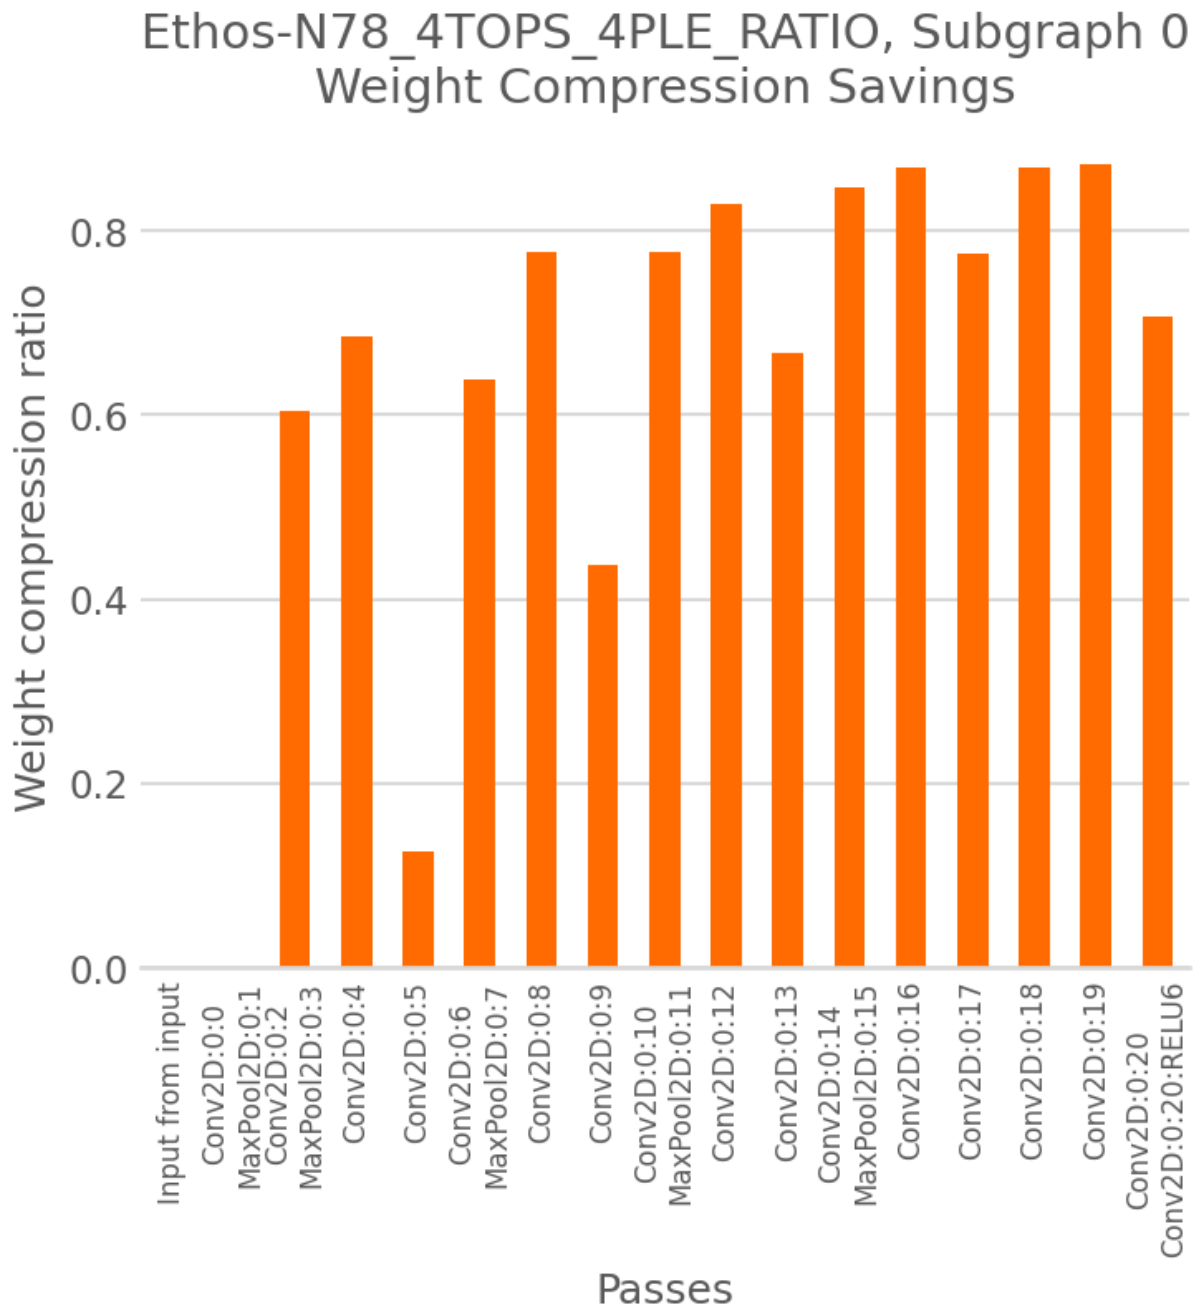
This is a breakdown of the weight compression of each pass

The weight compression savings for a pass is the percentage reduction of the space taken up by the weights, that has been saved by using compression. Higher values therefore indicate more compression has taken place.

The value is a "Space Saving" ratio, following the formula: $1 - (\text{compressed} / \text{uncompressed})$

```
In [11]: for subgraph, (_, data, _) in enumerate(results):
          df = spa.extract_weight_compression(data)
          ax = df.plot.bar(color=spa.ARM_COLORS['orange'])
```

```
'{0}', Subgraph {0}\nWeight Compression Savings'.format(params['Variant'], s)
spa.apply_formatting(ax, '{0}', Subgraph {0}\nWeight Compression Savings'.format
```



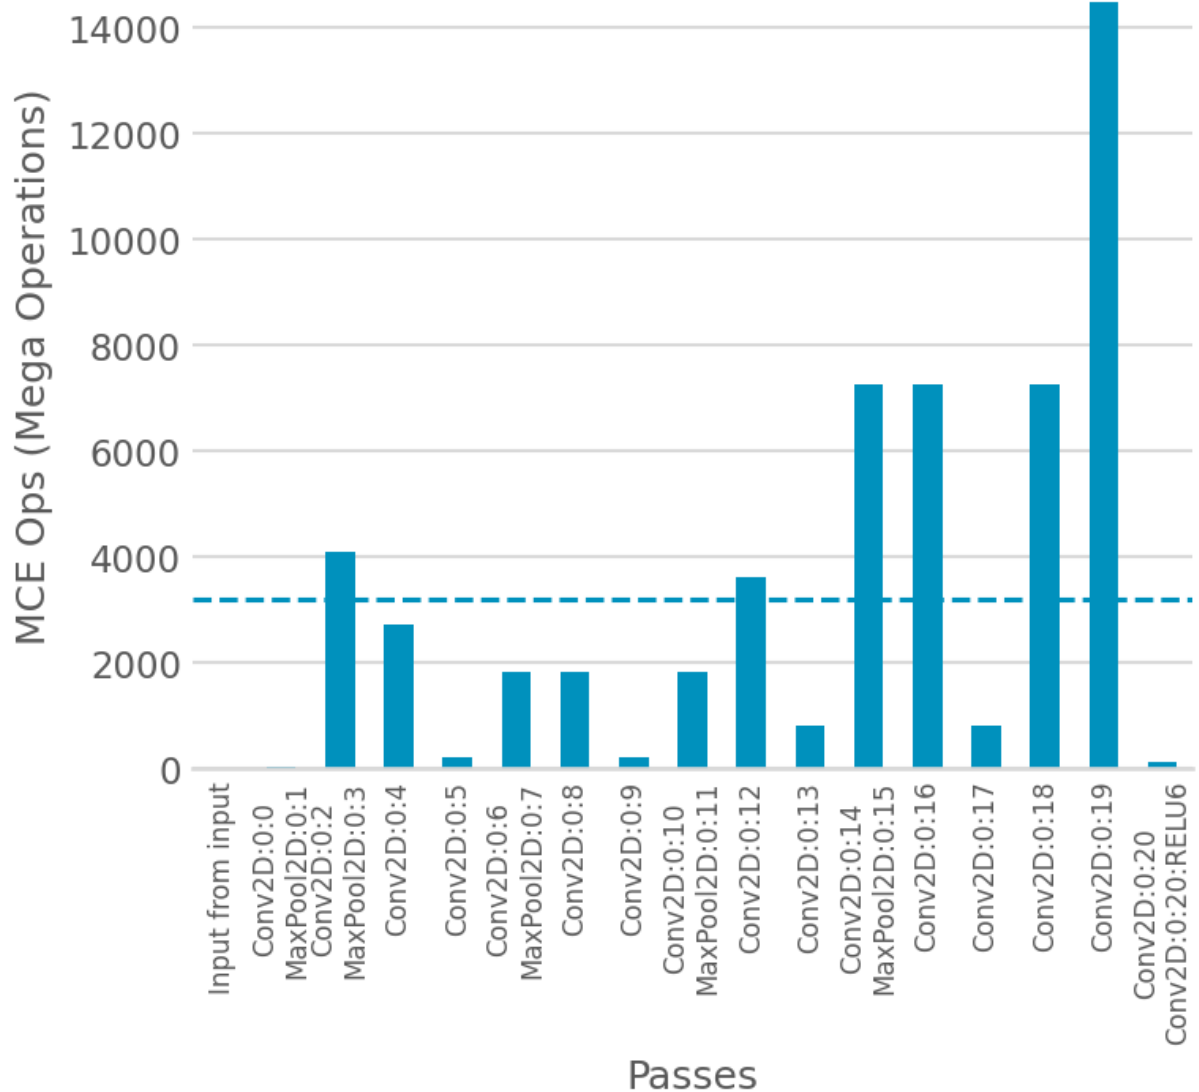
Network Complexity

This is a breakdown of complexity of each pass.

Complexity is the number of operations required for the MCE to perform that pass.

```
In [12]: for subgraph, (_, data, _) in enumerate(results):
df = spa.extract_mce_complexity(data, scale=1e6)
ax = df.plot.bar()
spa.apply_mean(ax, df)
spa.apply_formatting(ax, '{0}', Subgraph {0}\nNetwork Complexity (Ops)'.form
```

Ethos-N78_4TOPS_4PLE_RATIO, Subgraph 0 Network Complexity (Ops)



Cycle Count

This is a breakdown of cycles per pass.

The stacked bar chart displays a cycle count for each major unit: MCE, PLE, DMA Read, DMA Write; as well as the current firmware overhead.

The line chart displays the total predicted number of cycles it took to perform that pass.

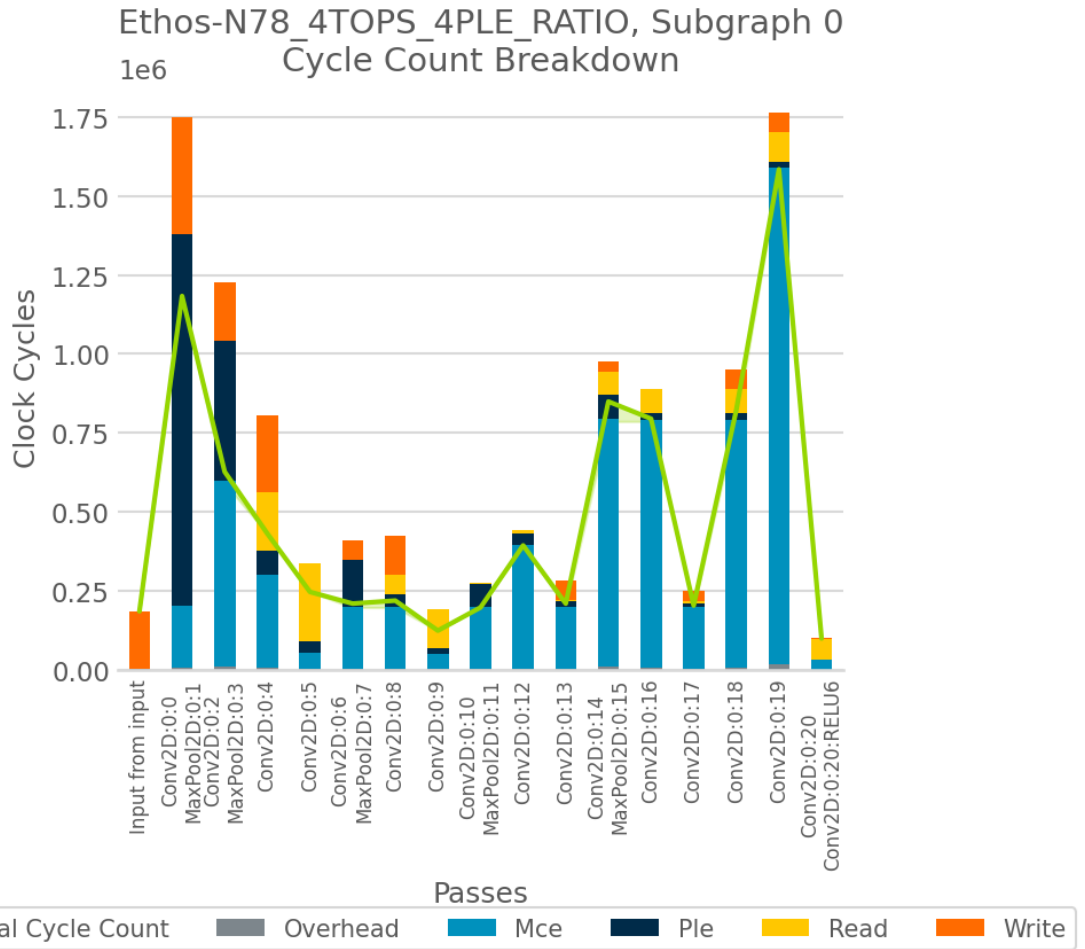
This line can be lower than the stacked chart, and when this happens, it denotes that parallelism is occurring.

The line also has a shaded area around it, which denotes theoretical upper and lower bounds for the predicted cycle count.

```
In [13]: for subgraph, (_, data, _) in enumerate(results):
          totals, df = spa.extract_cycle_count(data, BANDWIDTH, CLK)
          lower = totals.xs('Lower', axis=1, level=1).squeeze()
          upper = totals.xs('Upper', axis=1, level=1).squeeze()
          predicted = totals.xs('Predicted', axis=1, level=1).squeeze()
```

```

spa.plt.fill_between(totals.index, lower, upper, color=spa.ARM_COLORS['gr
ax = predicted.plot.line(color=spa.ARM_COLORS['green'], lw=2)
palette = [spa.ARM_COLORS['dark gray']] + spa.COOL_COLOR_PALETTE[1:] + sp
df.plot.bar(ax=ax, stacked=True, color=palette)
spa.apply_formatting(ax, '{}', Subgraph {} \nCycle Count Breakdown'.format(
    
```



Graph of operations

A graph representation of the operations that would be executed on the NPU.

Each node in this graph represents a command that would be passed to the hardware. Each command may be doing computations that correspond to more than one layer in the original network. Similarly a layer might be split across multiple commands.

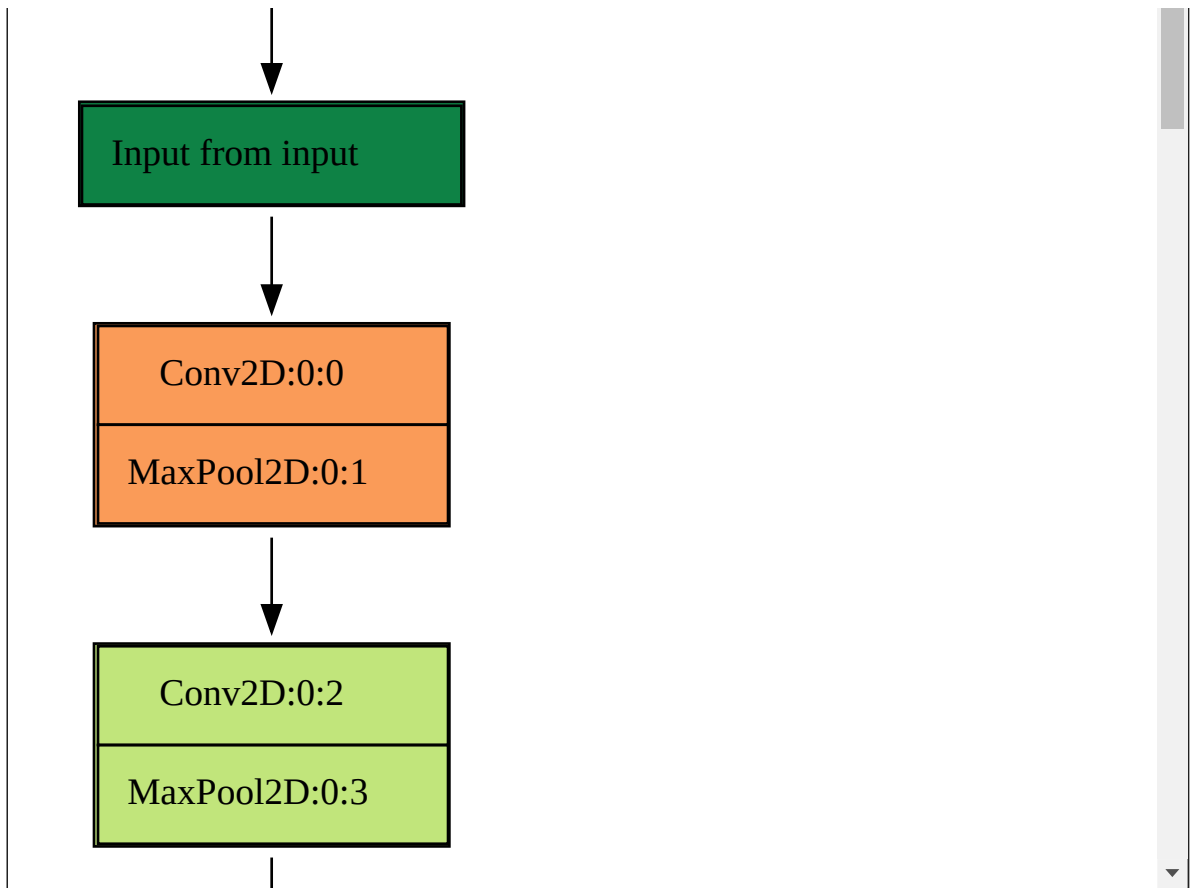
Nodes are color coded with a RAG rated color map based on the estimated number of cycles that they would take to run. Red for nodes that take longer, green for nodes that take less time.

```

In [14]: for subgraph, (_, data, _) in enumerate(results):
# Heatmap the graph using the cycle count
cycles = spa.calc_total_cycles_range(data, BANDWIDTH / CLK)['Predicted']
svg = spa.get_graph_svg(data, cycles)
print('Subgraph', subgraph)
html = '<div style="max-height:500px;border-style:solid;border-width:1px;
display(HTML(html))
    
```

Subgraph 0






Detailed Report

```
In [15]: for subgraph, (_, data, _) in enumerate(results):
# Display the parsed data in raw form
print('Subgraph', subgraph)
display(data)
```

Subgraph 0

	DramParallelBytes	DramNonParallelBytes	SramBytes	NumCentralStripes	NumBoundaryStripes
0	0.0	0.0	2359296.0	12.0	0.0
1	0.0	0.0	18874368.0	0.0	0.0
2	0.0	0.0	4915200.0	32.0	16.0
3	2162688.0	196608.0	0.0	24.0	0.0
4	2516582.0	629145.0	0.0	12.0	0.0
5	0.0	0.0	1572864.0	12.0	0.0

	DramParallelBytes	DramNonParallelBytes	SramBytes	NumCentralStripes	NumBoundaryStripes
6	655360.0	131072.0	0.0	12.0	0.0
7	1258291.0	314572.0	0.0	6.0	0.0
8	0.0	0.0	786432.0	6.0	0.0
9	0.0	0.0	393216.0	1.0	0.0
10	0.0	0.0	1572864.0	6.0	0.0
11	0.0	786432.0	0.0	1.0	0.0
12	314572.0	78643.0	0.0	1.0	0.0
13	0.0	0.0	786432.0	1.0	0.0
14	0.0	393216.0	0.0	1.0	0.0
15	0.0	0.0	786432.0	1.0	0.0
16	0.0	786432.0	0.0	1.0	0.0



In []:

In []: